

Amardeep Singh

E23CSEU2189

```
!wget http://efrosgans.eecs.berkeley.edu/pix2pix/datasets/edges2shoe
!tar -xvf edges2shoes.tar.gz
```

```
edges2shoes/train/47201_AB.jpg
edges2shoes/train/47282_AB.jpg
edges2shoes/train/47283_AB.jpg
edges2shoes/train/47284_AB.jpg
edges2shoes/train/47285_AB.jpg
edges2shoes/train/47286_AB.jpg
edges2shoes/train/47287_AB.jpg
edges2shoes/train/47288_AB.jpg
edges2shoes/train/47289_AB.jpg
edges2shoes/train/4728_AB.jpg
edges2shoes/train/47290_AB.jpg
edges2shoes/train/47291_AB.jpg
edges2shoes/train/47292_AB.jpg
edges2shoes/train/47293_AB.jpg
edges2shoes/train/47294_AB.jpg
edges2shoes/train/48843_AB.jpg
edges2shoes/train/48844_AB.jpg
edges2shoes/train/48845_AB.jpg
edges2shoes/train/48846_AB.jpg
edges2shoes/train/48847_AB.jpg
edges2shoes/train/48848_AB.jpg
edges2shoes/train/48849_AB.jpg
edges2shoes/train/4884_AB.jpg
edges2shoes/train/48850_AB.jpg
edges2shoes/train/48851_AB.jpg
edges2shoes/train/48852_AB.jpg
edges2shoes/train/48853_AB.jpg
edges2shoes/train/48854_AB.jpg
edges2shoes/train/48855_AB.jpg
edges2shoes/train/48856_AB.jpg
edges2shoes/train/48857_AB.jpg
edges2shoes/train/48858_AB.jpg
edges2shoes/train/48859_AB.jpg
edges2shoes/train/480_AB.jpg
edges2shoes/train/48100_AB.jpg
edges2shoes/train/48101_AB.jpg
edges2shoes/train/48102_AB.jpg
edges2shoes/train/48103_AB.jpg
edges2shoes/train/48104_AB.jpg
edges2shoes/train/48105_AB.jpg
edges2shoes/train/48106_AB.jpg
edges2shoes/train/48107_AB.jpg
```

```
edges2shoes/train/48108_AB.jpg  
edges2shoes/train/48109_AB.jpg  
edges2shoes/train/4810_AB.jpg  
edges2shoes/train/48110_AB.jpg  
edges2shoes/train/48111_AB.jpg  
edges2shoes/train/48112_AB.jpg  
edges2shoes/train/48113_AB.jpg  
edges2shoes/train/48114_AB.jpg  
edges2shoes/train/48115_AB.jpg  
edges2shoes/train/48116_AB.jpg  
edges2shoes/train/48117_AB.jpg  
edges2shoes/train/48118_AB.jpg  
edges2shoes/train/48119_AB.jpg  
edges2shoes/train/49521_AB.jpg  
edges2shoes/train/49522_AB.jpg  
edges2shoes/train/49523_AB.jpg  
edges2shoes/train/49524_AB.jpg
```

```
import os  
from PIL import Image  
  
os.makedirs("dataset/train/input", exist_ok=True)  
os.makedirs("dataset/train/target", exist_ok=True)  
  
source_folder = "edges2shoes/train"  
  
for i, img_name in enumerate(os.listdir(source_folder)):  
    if i > 1000: # LIMIT for faster training (remove this line for  
        break  
  
    img = Image.open(os.path.join(source_folder, img_name))  
    w, h = img.size  
  
    input_img = img.crop((0, 0, w//2, h))  
    target_img = img.crop((w//2, 0, w, h))  
  
    input_img.save(f"dataset/train/input/{img_name}")  
    target_img.save(f"dataset/train/target/{img_name}")  
  
print("Dataset Prepared Successfully!")
```

Dataset Prepared Successfully!

```
import torch  
import torch.nn as nn  
import torch.optim as optim  
from torchvision import transforms  
from torch.utils.data import Dataset, DataLoader  
from PIL import Image
```

```
import os
import matplotlib.pyplot as plt

# =====
# CONFIG
# =====
DEVICE = torch.device("cuda" if torch.cuda.is_available() else "cpu")
BATCH_SIZE = 4
LR = 0.0002
EPOCHS = 50
LAMBDA_L1 = 100
IMAGE_SIZE = 256

# =====
# DATASET
# =====
class PairedDataset(Dataset):
    def __init__(self, input_dir, target_dir, transform=None):
        self.input_dir = input_dir
        self.target_dir = target_dir
        self.transform = transform
        self.images = os.listdir(input_dir)

    def __len__(self):
        return len(self.images)

    def __getitem__(self, idx):
        img_name = self.images[idx]
        input_img = Image.open(os.path.join(self.input_dir, img_name))
        target_img = Image.open(os.path.join(self.target_dir, img_name))

        if self.transform:
            input_img = self.transform(input_img)
            target_img = self.transform(target_img)

        return input_img, target_img

    transform = transforms.Compose([
        transforms.Resize((IMAGE_SIZE, IMAGE_SIZE)),
        transforms.ToTensor(),
        transforms.Normalize((0.5,), (0.5,))
    ])

train_dataset = PairedDataset(
    "dataset/train/input",
    "dataset/train/target",
    transform
)
```

```
train_loader = DataLoader(train_dataset, batch_size=BATCH_SIZE, shu  
  
# =====  
# U-NET GENERATOR  
# =====  
class Down(nn.Module):  
    def __init__(self, in_c, out_c, normalize=True):  
        super().__init__()  
        layers = [nn.Conv2d(in_c, out_c, 4, 2, 1, bias=False)]  
        if normalize:  
            layers.append(nn.BatchNorm2d(out_c))  
        layers.append(nn.LeakyReLU(0.2))  
        self.model = nn.Sequential(*layers)  
  
    def forward(self, x):  
        return self.model(x)  
  
class Up(nn.Module):  
    def __init__(self, in_c, out_c):  
        super().__init__()  
        self.model = nn.Sequential(  
            nn.ConvTranspose2d(in_c, out_c, 4, 2, 1, bias=False),  
            nn.BatchNorm2d(out_c),  
            nn.ReLU(True)  
        )  
  
    def forward(self, x, skip):  
        x = self.model(x)  
        return torch.cat([x, skip], 1)  
  
class Generator(nn.Module):  
    def __init__(self):  
        super().__init__()  
        self.d1 = Down(3, 64, False)  
        self.d2 = Down(64, 128)  
        self.d3 = Down(128, 256)  
        self.d4 = Down(256, 512)  
  
        self.u1 = Up(512, 256)  
        self.u2 = Up(512, 128)  
        self.u3 = Up(256, 64)  
  
        self.final = nn.Sequential(  
            nn.ConvTranspose2d(128, 3, 4, 2, 1),  
            nn.Tanh()  
        )
```

```
def forward(self, x):
    d1 = self.d1(x)
    d2 = self.d2(d1)
    d3 = self.d3(d2)
    d4 = self.d4(d3)

    u1 = self.u1(d4, d3)
    u2 = self.u2(u1, d2)
    u3 = self.u3(u2, d1)

    return self.final(u3)

# =====
# PATCHGAN DISCRIMINATOR
# =====
class Discriminator(nn.Module):
    def __init__(self):
        super().__init__()
        self.model = nn.Sequential(
            nn.Conv2d(6, 64, 4, 2, 1),
            nn.LeakyReLU(0.2),

            nn.Conv2d(64, 128, 4, 2, 1),
            nn.BatchNorm2d(128),
            nn.LeakyReLU(0.2),

            nn.Conv2d(128, 256, 4, 2, 1),
            nn.BatchNorm2d(256),
            nn.LeakyReLU(0.2),

            nn.Conv2d(256, 1, 4, 1, 1)
        )

    def forward(self, input_img, target_img):
        x = torch.cat([input_img, target_img], dim=1)
        return self.model(x)

# =====
# BASELINE CNN
# =====
class BaselineCNN(nn.Module):
    def __init__(self):
        super().__init__()
        self.encoder = nn.Sequential(
            nn.Conv2d(3, 64, 4, 2, 1),
            nn.ReLU(),
            nn.Conv2d(64, 128, 4, 2, 1),
            nn.ReLU()
```

```
)  
  
        self.decoder = nn.Sequential(  
            nn.ConvTranspose2d(128, 64, 4, 2, 1),  
            nn.ReLU(),  
            nn.ConvTranspose2d(64, 3, 4, 2, 1),  
            nn.Tanh()  
    )  
  
    def forward(self, x):  
        return self.decoder(self.encoder(x))  
  
# =====  
# INITIALIZE MODELS  
# =====  
generator = Generator().to(DEVICE)  
discriminator = Discriminator().to(DEVICE)  
baseline = BaselineCNN().to(DEVICE)  
  
opt_G = optim.Adam(generator.parameters(), lr=LR, betas=(0.5, 0.999))  
opt_D = optim.Adam(discriminator.parameters(), lr=LR, betas=(0.5, 0.999))  
opt_base = optim.Adam(baseline.parameters(), lr=LR)  
  
criterion_GAN = nn.MSELoss()  
criterion_L1 = nn.L1Loss()  
  
# =====  
# TRAINING LOOP  
# =====  
for epoch in range(EPOCHS):  
    for input_img, target_img in train_loader:  
        input_img = input_img.to(DEVICE)  
        target_img = target_img.to(DEVICE)  
  
        # -----  
        # Train Discriminator  
        # -----  
        opt_D.zero_grad()  
        fake_img = generator(input_img)  
  
        real_pred = discriminator(input_img, target_img)  
        fake_pred = discriminator(input_img, fake_img.detach())  
  
        loss_real = criterion_GAN(real_pred, torch.ones_like(real_pred))  
        loss_fake = criterion_GAN(fake_pred, torch.zeros_like(fake_pred))  
        loss_D = (loss_real + loss_fake) * 0.5  
  
        loss_D.backward()
```

```
opt_D.step()

# -----
# Train Generator
# -----
opt_G.zero_grad()
fake_pred = discriminator(input_img, fake_img)

loss_GAN = criterion_GAN(fake_pred, torch.ones_like(fake_pr
loss_L1 = criterion_L1(fake_img, target_img) * LAMBDA_L1
loss_G = loss_GAN + loss_L1

loss_G.backward()
opt_G.step()

# -----
# Train Baseline CNN
# -----
opt_base.zero_grad()
base_out = baseline(input_img)
loss_base = criterion_L1(base_out, target_img)
loss_base.backward()
opt_base.step()

print(f"Epoch [{epoch+1}/{EPOCHS}] | D: {loss_D.item():.4f} | G: {loss_G.item():.4f} | Base: {loss_base.item():.4f}")

# =====
# SAVE MODELS
# =====
torch.save(generator.state_dict(), "generator.pth")
torch.save(discriminator.state_dict(), "discriminator.pth")
torch.save(baseline.state_dict(), "baseline.pth")

print("Training Complete!")
```

Epoch	[{}/50]	D:	G:	Base:
Epoch	[1/50]	0.1326	23.8385	0.3046
Epoch	[2/50]	0.1740	23.6484	0.4318
Epoch	[3/50]	0.2696	20.5333	0.4286
Epoch	[4/50]	0.2144	16.2506	0.5296
Epoch	[5/50]	0.0709	12.5423	0.2945
Epoch	[6/50]	0.0793	23.4040	0.2501
Epoch	[7/50]	0.0875	28.0215	0.2757
Epoch	[8/50]	0.1936	12.7036	0.3238
Epoch	[9/50]	0.3700	17.4188	0.4683
Epoch	[10/50]	0.0661	29.0178	0.3511
Epoch	[11/50]	0.1447	12.4998	0.5132
Epoch	[12/50]	0.0819	13.6663	0.3779
Epoch	[13/50]	0.0919	14.3000	0.3940
Epoch	[14/50]	0.0555	12.9992	0.4653

Epoch [15/50]	D: 0.3715	G: 13.9545	Base: 0.4313
Epoch [16/50]	D: 0.2921	G: 8.2361	Base: 0.1598
Epoch [17/50]	D: 0.1919	G: 13.8081	Base: 0.3025
Epoch [18/50]	D: 0.1326	G: 13.0462	Base: 0.6527
Epoch [19/50]	D: 0.2411	G: 5.2493	Base: 0.2288
Epoch [20/50]	D: 0.0493	G: 9.2747	Base: 0.3478
Epoch [21/50]	D: 0.2254	G: 11.3301	Base: 0.6343
Epoch [22/50]	D: 0.2876	G: 15.3689	Base: 0.3124
Epoch [23/50]	D: 0.1354	G: 6.4848	Base: 0.2591
Epoch [24/50]	D: 0.0393	G: 5.6221	Base: 0.4948
Epoch [25/50]	D: 0.1374	G: 13.6044	Base: 0.2861
Epoch [26/50]	D: 0.1897	G: 3.9318	Base: 0.2597
Epoch [27/50]	D: 0.1872	G: 9.1262	Base: 0.5169
Epoch [28/50]	D: 0.1604	G: 9.2063	Base: 0.4226
Epoch [29/50]	D: 0.2314	G: 7.1753	Base: 0.2388
Epoch [30/50]	D: 0.0972	G: 8.4770	Base: 0.4747
Epoch [31/50]	D: 0.1922	G: 6.6966	Base: 0.0946
Epoch [32/50]	D: 0.2732	G: 6.1737	Base: 0.2416
Epoch [33/50]	D: 0.1113	G: 6.8787	Base: 0.1766
Epoch [34/50]	D: 0.2740	G: 3.8584	Base: 0.5602
Epoch [35/50]	D: 0.1239	G: 16.1348	Base: 0.4111
Epoch [36/50]	D: 0.2689	G: 9.2613	Base: 0.4196
Epoch [37/50]	D: 0.1940	G: 6.5232	Base: 0.6022
Epoch [38/50]	D: 0.0922	G: 5.6428	Base: 0.3318
Epoch [39/50]	D: 0.0779	G: 7.7947	Base: 0.3940
Epoch [40/50]	D: 0.1708	G: 5.6364	Base: 0.5595
Epoch [41/50]	D: 0.1839	G: 7.6923	Base: 0.3897
Epoch [42/50]	D: 0.2386	G: 12.7532	Base: 0.3503
Epoch [43/50]	D: 0.1116	G: 4.4683	Base: 0.3012
Epoch [44/50]	D: 0.3022	G: 4.1060	Base: 0.1722
Epoch [45/50]	D: 0.3117	G: 4.3732	Base: 0.2126
Epoch [46/50]	D: 0.2686	G: 7.9298	Base: 0.3121
Epoch [47/50]	D: 0.1358	G: 7.9810	Base: 0.3704
Epoch [48/50]	D: 0.1213	G: 9.1755	Base: 0.2516
Epoch [49/50]	D: 0.2128	G: 3.6445	Base: 0.3876
Epoch [50/50]	D: 0.2295	G: 5.8088	Base: 0.3301

Training Complete!

```
import matplotlib.pyplot as plt

generator.eval()

input_img, target_img = next(iter(train_loader))
input_img = input_img.to(DEVICE)

with torch.no_grad():
    fake_img = generator(input_img)

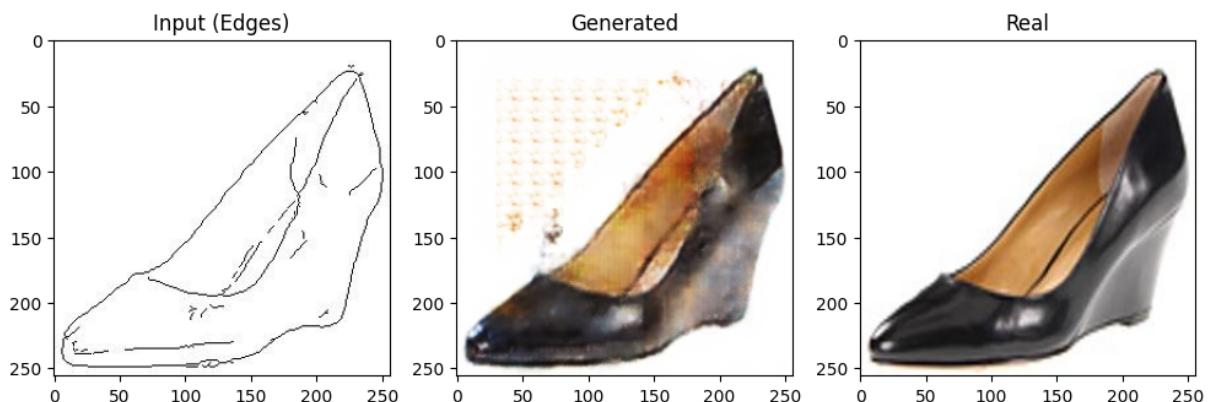
plt.figure(figsize=(12,4))

plt.subplot(1,3,1)
plt.title("Input (Edges)")
plt.imshow(input_img[0].permute(1,2,0).cpu()*0.5+0.5)

plt.subplot(1,3,2)
plt.title("Generated")
plt.imshow(fake_img[0].permute(1,2,0).cpu()*0.5+0.5)

plt.subplot(1,3,3)
plt.title("Real")
plt.imshow(target_img[0].permute(1,2,0)*0.5+0.5)

plt.show()
```



```
import os
os.makedirs("generated_outputs", exist_ok=True)
print("Folder created:", os.path.exists("generated_outputs"))
```

Folder created: True

```
import torch
import matplotlib.pyplot as plt
import os

generator.eval()

count = 0

for input_img, target_img in train_loader:
    input_img = input_img.to(DEVICE)

    with torch.no_grad():
        fake_img = generator(input_img)

    for i in range(fake_img.size(0)):
        fake = fake_img[i].permute(1,2,0).cpu()*0.5+0.5
        fake = fake.clamp(0,1)

        plt.imsave(f"generated_outputs/output_{count}.png", fake)
        count += 1

    if count >= 10:
        break
if count >= 10:
    break

print("Saved files:", os.listdir("generated_outputs"))
```

Saved files: ['output_5.png', 'output_3.png', 'output_7.png', 'output_1.png']

```
!zip -r outputs.zip generated_outputs
```

```
adding: generated_outputs/ (stored 0%)
adding: generated_outputs/output_5.png (deflated 0%)
adding: generated_outputs/output_3.png (deflated 0%)
adding: generated_outputs/output_7.png (deflated 0%)
adding: generated_outputs/output_1.png (deflated 0%)
adding: generated_outputs/output_8.png (deflated 0%)
adding: generated_outputs/output_0.png (deflated 0%)
adding: generated_outputs/output_2.png (deflated 0%)
adding: generated_outputs/output_6.png (deflated 0%)
adding: generated_outputs/output_9.png (deflated 0%)
adding: generated_outputs/output_4.png (deflated 0%)
```

```
from google.colab import files
files.download("outputs.zip")
```